

Assignement 9

1 Training an Autoencoder as a 2D Feature Generator and Visualizing CIFAR-10 Features

The objective of this experiment is to train a convolutional autoencoder on the CIFAR-10 dataset in an unsupervised manner. The encoder is designed to learn compressed 2-dimensional representations (features) of input images. These 2D features are later visualized using scatter plots to observe how different classes are separated in the latent space.

1.1 Dataset

- **Name:** CIFAR-10
- **Size:** 60,000 images ($32 \times 32 \times 3$)
 - 50,000 for training
 - 10,000 for testing
- **Classes (10 total):**
Airplane, Automobile, Bird, Cat, Deer, Dog, Frog, Horse, Ship, Truck

1.2 Methodology

1.2.1 Preprocessing

- Images are normalized to the range $[0, 1]$.
- All 60,000 images (train + test) are combined for unsupervised training.
- Corresponding class labels are concatenated to color points in the feature space.

1.2.2 Model Architecture

Encoder:

- Conv2D(32) \rightarrow MaxPooling
- Conv2D(64) \rightarrow MaxPooling
- Flatten \rightarrow Dense(2)

Decoder:

- Dense \rightarrow Reshape
- Conv2DTranspose layers to reconstruct the input

1.2.3 Training

- **Loss Function:** Mean Squared Error (MSE)
- **Optimizer:** Adam
- **Epochs:** 10
- **Batch Size:** 256
- **Data:** All 60,000 images (unsupervised)

1.3 Visualization

A scatter plot is created from the 2D features extracted by the encoder. Points are color-coded according to their class labels using `matplotlib` and the `tab10` colormap.

1.4 Result

- The 2D encoded features show **clear clustering** based on semantic similarity.
- Classes such as *Airplane*, *Truck*, and *Ship* show visible separation, while visually similar classes like *Cat* and *Dog* may overlap.
- The autoencoder successfully learns compressed representations that **preserve class-level structure** in an unsupervised manner.




Figure 1: 2D Encoded Feature Visualization of CIFAR-10

2 Comparing autoencoder generated features with features extracted by a pre-trained CNN and reduced by dimension reduction techniques like PCA, t-SNE.

2.1 Dataset

- **Dataset:** CIFAR-10
- **Number of Images:** 5000 (random subset)
- **Image Size:** 32×32 RGB
- **Classes:** 10

2.2 Implementation Steps

1. Autoencoder Training:

- Input shape: (32, 32, 3)
- Encoder compresses features to 2D (bottleneck layer)
- Decoder reconstructs images from the 2D features
- Trained using Adam optimizer and Mean Squared Error (MSE) loss for 10 epochs

2. Feature Extraction using Autoencoder:

- Extracted 2D encoded features
- Directly visualized using scatter plot with class labels

3. Feature Extraction using Pretrained VGG16:

- Images resized to 224×224×3
- Preprocessed using VGG16-specific normalization
- Features extracted from the global average pooling layer (last layer before classification)

4. Dimensionality Reduction:

- PCA applied to reduce VGG16 features from 512-D to 2-D
- t-SNE used for nonlinear 2D projection

5. Visualization:

- Used scatter plots for all 2D projections
- Points color-coded by CIFAR-10 class

Results (Visualizations)




Figure 2: Autoencoder 2D Features




Figure 3: VGG16 Features + PCA

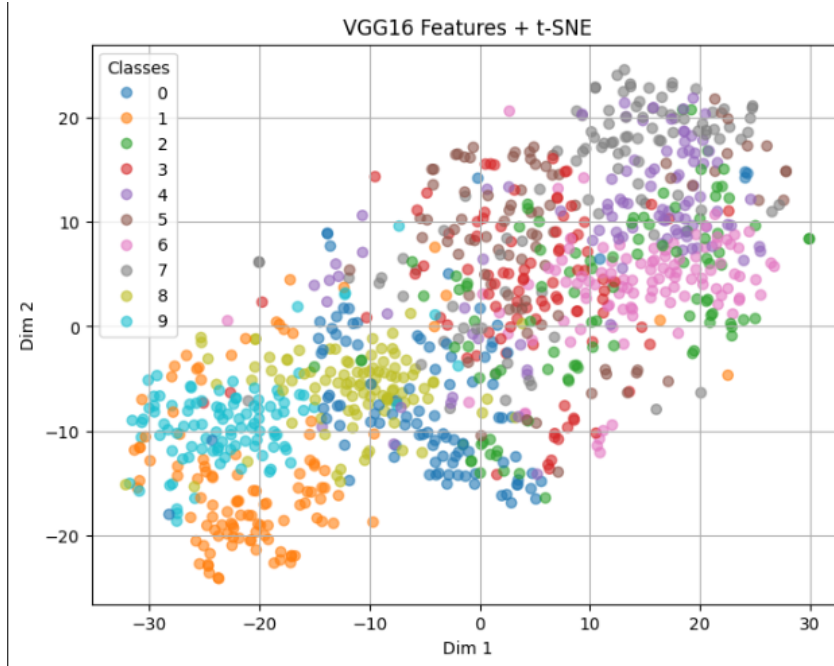


Figure 4: VGG16 Features + t-SNE

3 Training a Denoising Autoencoder on the CIFAR-10 Dataset

3.1 Dataset

- **Dataset Used:** CIFAR-10
- **Image Size:** $32 \times 32 \times 3$
- **Training Samples:** 50,000
- **Testing Samples:** 10,000
- **Preprocessing:**
 - Pixel values normalized to the $[0, 1]$ range
 - Gaussian noise (mean = 0, std = 1, factor = 0.2) added to training and test data
 - Noisy values clipped to $[0, 1]$

3.2 Model Architecture

3.2.1 Encoder

- Conv2D (32 filters, 3×3 , ReLU, padding='same')
- MaxPooling2D (2×2 , padding='same')
- Conv2D (64 filters, 3×3 , ReLU, padding='same')
- MaxPooling2D (2×2 , padding='same')

3.2.2 Decoder

- Conv2D (64 filters, 3×3 , ReLU, padding='same')
- UpSampling2D (2×2)
- Conv2D (32 filters, 3×3 , ReLU, padding='same')
- UpSampling2D (2×2)
- Conv2D (3 filters, 3×3 , Sigmoid, padding='same') — output layer

3.3 Training Details

- **Loss Function:** Binary Crossentropy
- **Optimizer:** Adam
- **Epochs:** 10
- **Batch Size:** 128
- **Validation:** Performed on noisy test images with their original versions as targets

3.4 Results(Visualization)

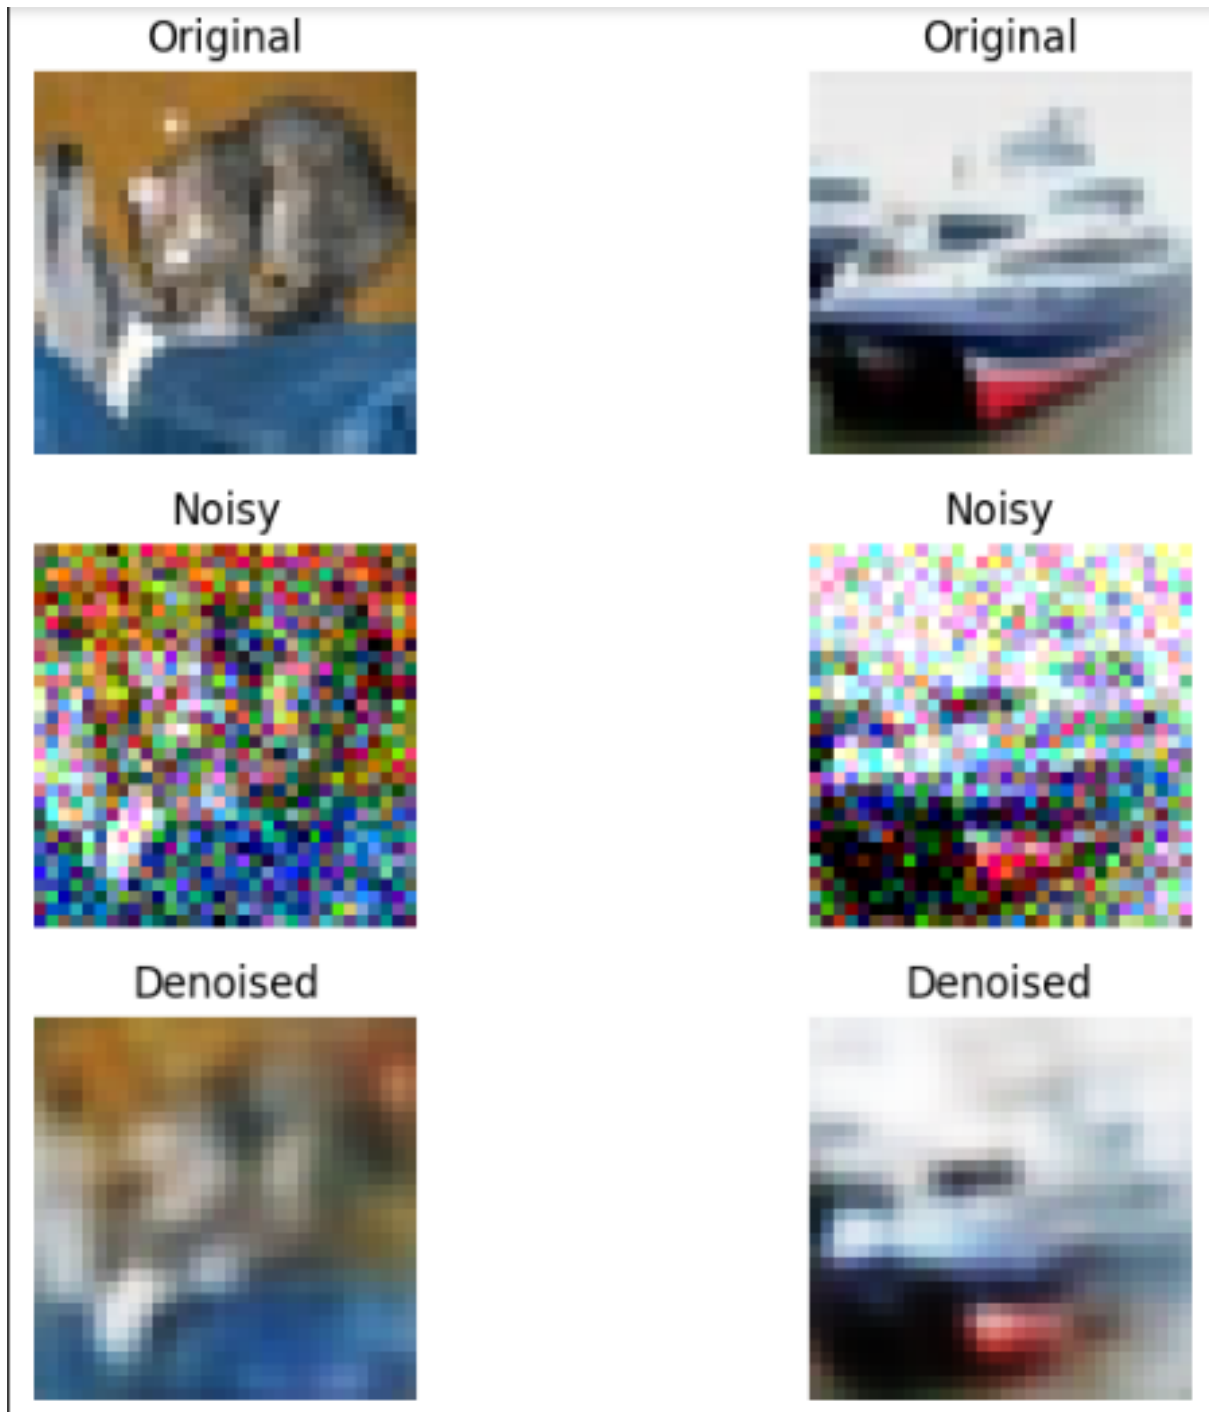


Figure 5: Example of Original, Noisy, and Denoised images from the CIFAR-10 dataset.

4 Training a CNN based CIFAR-10 classifier without any single-image data augmentation techniques

4.1 Dataset

- **CIFAR-10:** Contains 60,000 32×32 color images in 10 classes, with 6,000 images per class.

- **Split:** 50,000 training images and 10,000 test images.

4.2 Preprocessing

- Image pixel values normalized to the range $[0, 1]$.
- Labels were one-hot encoded for multi-class classification.

4.3 Model Architecture

- **Input:** $32 \times 32 \times 3$
- **Convolutional Block 1:**
 - Conv2D(32, 3×3 , ReLU, same padding)
 - Conv2D(32, 3×3 , ReLU)
 - MaxPooling2D(2×2)
 - Dropout(0.25)
- **Convolutional Block 2:**
 - Conv2D(64, 3×3 , ReLU, same padding)
 - Conv2D(64, 3×3 , ReLU)
 - MaxPooling2D(2×2)
 - Dropout(0.25)
- **Fully Connected:**
 - Flatten
 - Dense(512, ReLU)
 - Dropout(0.5)
 - Dense(10, Softmax)

4.4 Training Configuration

- **Loss Function:** Categorical Crossentropy
- **Optimizer:** Adam
- **Metrics:** Accuracy
- **Epochs:** 20
- **Batch Size:** 64
- **Validation:** Test set used for validation

4.5 Results

The model was trained and validated over 20 epochs. Accuracy and loss curves were plotted to monitor performance.

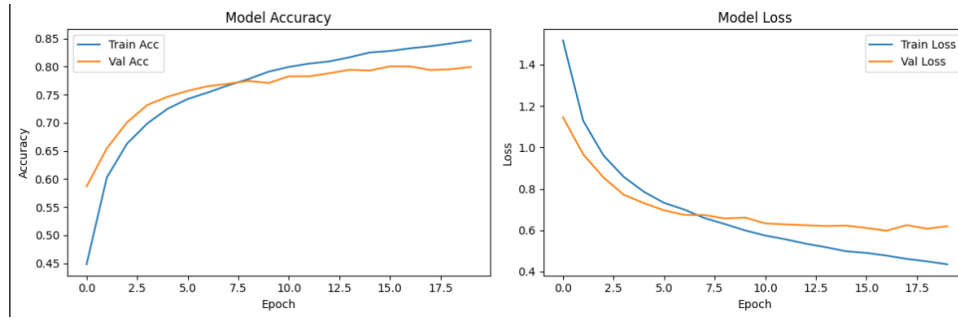


Figure 6: Training and validation accuracy (left) and loss (right) over epochs.

5 training a CNN based CIFAR-10 classifier with a single/multiple single-image data augmentation techniques

5.1 Dataset

- **Name:** CIFAR-10
- **Images:** 60,000 (32x32 RGB)
 - 50,000 training images
 - 10,000 test images
- **Classes:** 10 (airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck)

5.2 Data Preprocessing

- Normalized pixel values to the range $[0, 1]$.
- One-hot encoding applied to the labels.

5.3 Data Augmentation Techniques Applied

- Horizontal Flip
- Random Rotation ($\pm 10\%$)
- Random Zoom ($\pm 10\%$)
- Random Translation ($\pm 10\%$)
- Random Contrast ($\pm 10\%$)

5.4 Model Architecture

A CNN consisting of:

- Two convolutional blocks:

- Conv2D → Conv2D → MaxPooling2D → Dropout
- Fully connected layers:
 - Flatten → Dense(512) → Dropout → Dense(10)

5.5 Training Details

- **Epochs:** 20
- **Batch size:** 64
- **Loss function:** Categorical Crossentropy
- **Optimizer:** Adam
- **Metrics:** Accuracy

5.6 Results

5.6.1 Model Accuracy and Loss Curves




Figure 7: Training and Validation Accuracy and Loss

5.6.2 Visualization of Data Augmentation

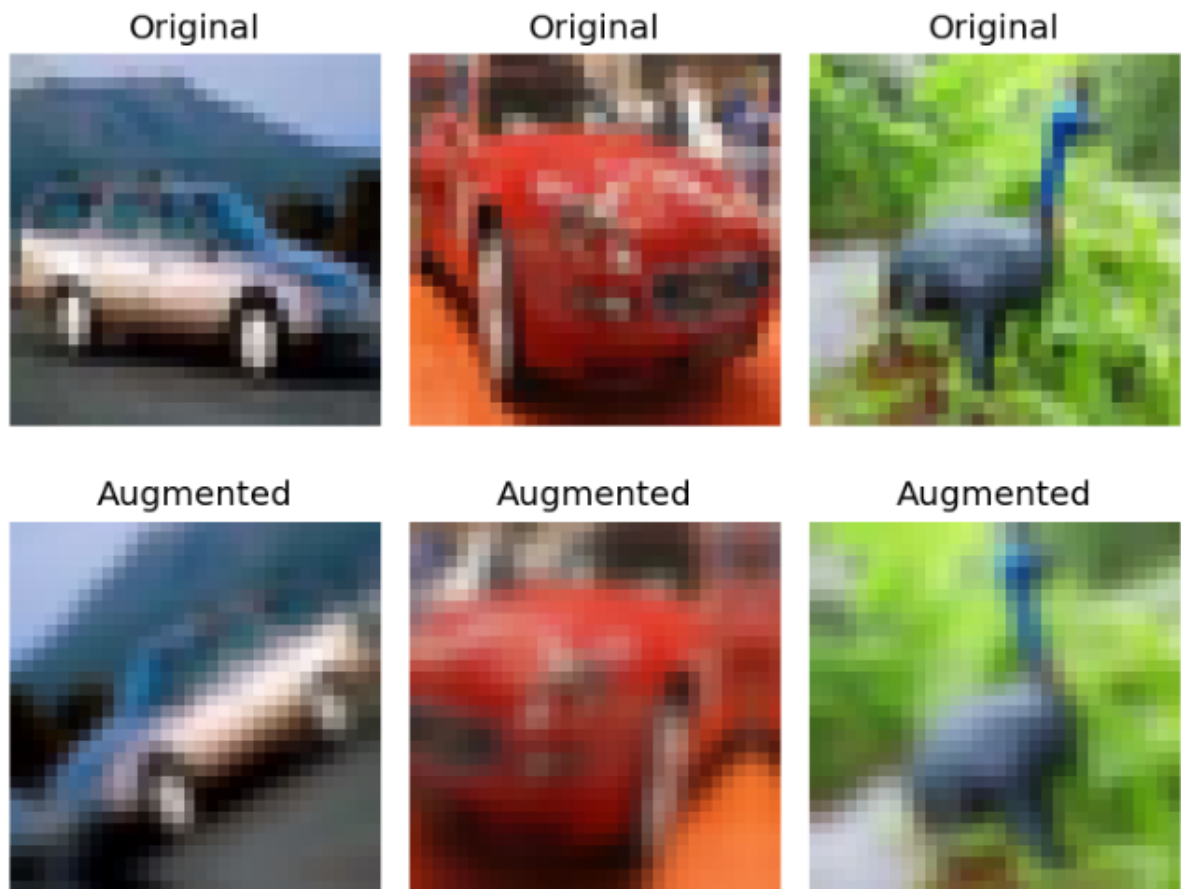


Figure 8: Examples of Augmented Images

Source Link: You can access the code and visuals from the following link:
Github Link